More Than Meets the Evidence: A New Look at Reasons for Belief

H. Brigitte Gill

hbgill@utexas.edu

Abstract

Many have thought that only evidential considerations can be reasons for which we believe. This is in large part because we've assumed doxastic transparency is true: that when we deliberate about what to believe, the question of whether to believe that p gives way to the question of whether p is true. This paper challenges those assumptions by drawing attention to the different functional roles reasons can play in deliberation. In different roles, reasons engage different deliberative questions. Not all of these neatly collapse into that of whether p is true: transparency is false. The illusion to the contrary is generated when we focus on overly simplistic cases of deliberation that encourage us to neglect the existence of what I term 'indirect' reasons. The implications for popular arguments about normative evidentialism and the aim or norm of belief are grim.

Key Words: transparency, doxastic deliberation, reasons, motivating reasons, exclusivity, evidentialism

1 Introduction

Imagine the weatherman forecasts rain for tomorrow. His word is perfectly good evidence of rain, so you believe it will rain on the basis of his testimony. That seems straightforward enough. Now, consider a different scenario: you're undecided about tomorrow's weather with no evidence either way. The weatherman offers you $100 if you can convince yourself it will rain. Could you conclude deliberation in belief for the reason that you want the money? Presumably not. In the literature, observations like this are typically explained by appeal to doxastic transparency. According to transparency, when we deliberate about what to believe, the question of whether to believe that p collapses into the question of whether p is true. As a consequence, the only considerations we can see as relevant to our deliberation are those we take to bear on p's truth – like the weatherman's testimony, and unlike our desire for money.

Though transparency is a claim about the psychology of doxastic deliberation, it's often thought to have far-reaching normative implications for the ethics of belief. It features in many arguments about the aim or norm of belief, all of which are premised on the idea that transparency is a striking phenomenon that calls out for explanation (Archer, 2017, Owens, 2003, Shah, 2003, Steglich-Petersen, 2006, 2009, Sullivan-Bisset & Noordhof 2013, 2020, Sullivan-Bisset, 2018, Sylvan, 2016, Velleman, 2000). Many see it as an important limitation on belief-control (Bennett, 1990, Hieronymi, 2006, Winters, 1979, Williams 1973), which could make the responsibility we have for belief very different from the distinctively moral responsibility we can have for bodily actions. And importantly, it plays a pivotal role in Shah's influential argument for evidentialism (Shah, 2006).

Transparency is widely regarded as uncontroversial,[1] and is typically taken to entail that only evidential considerations – those we see as connected to p's truth – can be reasons for which we believe. This is odd, given that non-academic discourse tends to take for granted that reasons for belief can be non-evidential. Notice how common it is for ordinary people to assume that (perceived) moral imperatives to believe women are among the reasons for which they and others believe women about sexual assault. Similar discourse can be found on a variety of topics, including sexual harassment, gender identity, disability, race, and criminal injustice. Such a stark difference between the contemporary philosophical debate and ordinary folk discourse ought to give us pause. At the very least, it deserves an explanation.

Rather than dismissing the non-academic discourse as imprecise, I argue that we should reconsider transparency. The literature has tended to focus on oversimplified cases of deliberation, which has limited discussion and generated the illusion that doxastic deliberation is transparent. When we turn our attention to more complex cases, it is clear that there is more than one functional role for reasons to play in deliberation. In different roles, reasons engage different deliberative questions. Not all of these questions neatly collapse into that of whether p is true. Transparency is false.

To start, I'll provide background and clarify key terms. In (3), I'll present my argument, discuss the distinction between direct and indirect reasons, and explain why this distinction challenges transparency. I'll address objections in section (4).

2 Background

In this paper I focus on doxastic deliberation and motivating doxastic reasons. Doxastic deliberation, or reasoning about what to believe, is one among many ways that we can acquire beliefs. Deliberation does not always conclude – it may be interrupted, or you may grow tired of it – but when it does, it concludes in a doxastic attitude. Doxastic reasons are reasons for having doxastic attitudes, such as belief, disbelief, suspension of judgment, or increasing/decreasing credal states. While a normative doxastic reason

[1] For recent exceptions, see McHugh (2015) and Sharadin (2016). McHugh's argument hinges on a controversial distinction between taking one's evidence to require versus merely permit belief, while Sharadin's hinges on a case of 'self-fulfilling prophecy' that may seem the exception rather than the norm when it comes to belief formation. My argument doesn't turn on McHugh's controversial distinction, and unlike Sharadin's argument, applies to ordinary cases of belief formation.

is a reason that objectively favors a doxastic attitude, and in that sense is an objectively good reason to believe or disbelieve (increase or decrease your credence, etc.), a motivating doxastic reason is a reason for which an individual takes a doxastic attitude. In other words, motivating reasons are psychological, as opposed to normative (Dancy 2000, 2018; Alvarez & Way 2024). While this much is standard in the literature, beyond it, the question of what exactly motivating reasons are becomes controversial.[2] I hope to sidestep this by assuming that any consideration which can be an explicit premise in an agent's doxastic deliberation can be a motivating doxastic reason for her. While not all motivating reasons are deliberative premises, all deliberative premises are capable of being motivating reasons. For my purposes, a premise in doxastic deliberation will be a consideration we take to bear on the question of what to believe and explicitly take into account during doxastic deliberation. It also guides us toward or away from a conclusion accordingly.

Now, why the focus on motivating doxastic reasons? Recall again transparency:

> DOXASTIC TRANSPARENCY: From the first-person perspective in doxastic deliberation, the question of whether to believe that p collapses into the question of whether p is true.

Deliberation is the process of trying to answer a question through conscious reasoning. In doxastic deliberation, the question we're trying to answer is that of whether to believe that p. (Note that this is different from the question of whether to intend to believe that p, of whether to induce belief that p, etc.) If transparency is true, then the belief-question we're trying to answer simply becomes the question of whether p is true. It seems to follow that only considerations we take to bear on the question we're trying to answer – which, according to transparency, is just the truth question – can consciously sway or guide our reasoning in the way characteristic of explicit deliberative premises. By definition, considerations we take to bear on the truth-question are evidential considerations (more on this below). Paired with the common thought that a consideration can only be a reason for which we believe if it's capable of being an explicit premise in reasoning about what to believe, this entails that only evidential considerations can be reasons for which we believe. In this way, transparency entails the following:

> PSYCHOLOGICAL EXCLUSIVITY: Only evidential considerations can be motivating doxastic reasons.[3]

Note that this differs from the normative thesis that only evidential considerations are good, proper, or normative reasons for belief. From here on, the term 'exclusivity' will refer to this psychological thesis. For the reasons detailed above, it's widely assumed that transparency entails exclusivity. This makes exclusivity a good litmus test for transparency: if exclusivity is false, then transparency must be too. Hence, the focus on explicit deliberative premises. If C can be an explicit premise in doxastic deliberation,

[2] For an overview, see Alvarez & Way 2024. For recent discussions of the controversies about motivating reasons, see Cunningham 2021; Engel-Hawbecker 2024; Gregory 2016; Lord & Sylvan 2019; Kiesewetter 2022; and Vahid 2022.

[3] There are stronger and weaker versions of exclusivity. The weakest version simply says that only evidential considerations can be reasons for belief or disbelief. The strongest version includes credence modification and suspension of judgment. This paper undermines both versions.

then it can be a motivating doxastic reason. If a non-evidential consideration can be a motivating doxastic reason, then exclusivity is false – which means that transparency is as well.

The term 'evidential' will also become important, later on. For the purposes of this paper, this is the definition I am working with:

> EVIDENTIAL CONSIDERATION: A consideration counts as evidential for some agent just in case she takes it to be connected to p's truth while deliberating about whether to believe p.

Notice that this is a subjective, psychological notion: the very same consideration may count as evidential for one agent and non-evidential for another, or for the same agent in a different deliberative context. Perhaps this isn't the way the term should normally be used (I have no theoretical commitments on the matter), but I use it this way out of charity for my opponent.[4]

3 Cases and Argument

The main argument of this paper can be broken down roughly as follows:

1. Evidential considerations that play the role of 'indirect reasons' in doxastic deliberation are motivating doxastic reasons.
2. Non-evidential considerations can play the role of 'indirect reasons' in doxastic deliberation.
3. If (1) and (2), then non-evidential considerations can be motivating doxastic reasons.
4. If non-evidential considerations can be motivating doxastic reasons, then exclusivity and transparency are false.
5. (from 1, 2, 3, 4) Exclusivity and transparency are false.

[4] If we didn't interpret 'evidential' in roughly this way, then exclusivity would be vulnerable to fairly obvious counterexamples, or at least would depend heavily on the correct theory of evidence. For one, we sometimes reason correctly to belief that p from considerations that support p's truth in a way that isn't obviously evidential, except in the loose sense of being connected to p's truth (ex: we might reason from an argument's structure to belief that it is valid).

Second, people often reason to belief that p from considerations that are objectively unconnected to p's truth. Billy Joe may reason from the fact that a presidential candidate has experience in business to belief that the candidate is skilled at foreign policy. Objectively, this consideration may not be evidence for p at all (e.g., if experience in business is unrelated to foreign policy skill). We could instead define 'evidential' such that q is an 'evidential consideration' relative to p for some agent just in case q raises the probability of p given the agent's other doxastic attitudes. But this won't help either – sloppy or emotional reasoning could lead Billy Joe to temporarily fail to appreciate the evidential relationships which hold between his own beliefs, causing him to see experience in business as connected to foreign policy skill even if he'd see in a calmer moment that his beliefs don't support this at all.

My opponent probably does not wish to deny the possibility of cases like these, though they would pose serious problems for her if 'evidential consideration' means anything other than, "a consideration the agent sees (at least momentarily) as connected to p's truth."

Here I should introduce the distinction between direct and indirect reasons. Direct reasons are typically what first come to mind when we think of motivating doxastic reasons – I conclude in belief that it's about to rain on the basis of the storm clouds looming overhead, or substantially decrease my credence that I-10 will be faster than back roads upon receiving news of an accident on the interstate. But not all motivating doxastic reasons are direct reasons. Consider the following case:

> WEATHER PRANKS: Leticia is deliberating about whether to believe it will rain tomorrow. She thinks to herself:
>
> "The local weatherman has said that it will rain, which has me leaning strongly toward concluding deliberation in belief. In fact, if I were to conclude deliberation right now, I would conclude in belief. But then again, the weatherman makes prank forecasts twice a month, doesn't he? This means I need to make an adjustment: while his testimony should still be worth something (he is still more reliable than chance), it shouldn't carry as much weight in my reasoning as it was initially. So for this reason I'll take his testimony less seriously than I was originally inclined. I still lean slightly toward belief, but my evidence no longer seems sufficient for it. I'll suspend judgment."

In this case, the weatherman's testimony and his status as a prankster play different functional roles in Leticia's deliberation. His testimony plays a role like those in the examples above; it seems to weigh directly in favor of incoming rain, and contributes a kind of independent motivational sway to her deliberation. As a consequence, the testimony's ability to guide her toward belief in this episode does not crucially depend on her actively attending to other considerations. Like the initial examples in this section, the weatherman's testimony would have swayed Leticia even if it had been the only consideration she consciously entertained.[5]

Unlike the testimony, the weatherman's status as a prankster does not seem to bear directly on the issue of incoming rain, and in that respect doesn't seem to be an independent reason for belief. If it had been the only consideration Leticia considered, it would not (or we can imagine a case where it would not) have carried any weight in her deliberation at all. It would have seemed irrelevant. It acquires relevance only via the weatherman's testimony, and sways her by bearing on the question of *how to weigh the weatherman's testimony*, rather than on the question of incoming rain. The fact that the weatherman is a prankster means his testimony should be worth less in her deliberation than it would have been otherwise.

Without the weatherman's testimony, his history of pulling pranks is not evidence for incoming rain. It does not bear on p's truth or falsity. Nor is it evidence for (or against) the reason it attenuates – that the weatherman is a casual prankster is not

[5]This is not the same as saying that the weatherman's testimony could have swayed Leticia completely independently of any background beliefs. Surely, she could not be moved by his testimony without the background belief that weathermen are generally reliable, that she was not hallucinating when she thought she heard the weatherman give his forecast, and so on. While her deliberation may have depended upon these background beliefs, she didn't consciously consider them in her reasoning, and so they were not explicit deliberative premises in this particular episode of deliberation – though they probably could be in other contexts.

evidence that he has or hasn't forecasted incoming rain. In this case, the sole function of the consideration that the weatherman pulls pranks is to attenuate the weight or role that his testimony plays in Leticia's reasoning. Reasons that function merely to attenuate other reasons is a well-documented phenomenon. Jonathan Dancy refers to reasons that perform this function in particular as 'attenuators/intensifiers' (Dancy, 2018).[6]

Something in the ballpark of this distinction is also widely recognized in the literature on defeaters. Rebutting defeaters for p present us with a reason for thinking that p is false, whereas undermining defeaters for p present us with a reason for thinking the usual connections between p and reasons for p do not hold. An undermining defeater may present you with evidence that belief that q is not justified, that q is not actually evidence for p, or that q does not support p to the degree you had originally thought (Pollock & Cruz, 1999, pp. 96-7). In other words, a rebutting defeater for p presents us with a reason for belief very much like the weatherman's testimony – a reason that directly favors p's truth or falsity. In contrast, an undermining defeater for p presents us with a reason like the weatherman's pranks, which bears instead on how other reasons are to be evaluated.

Of course, testimony does not always play a direct role, and factors pertaining to reliability are not always relegated to purely indirect roles. This is fine. The point of the example is simply that there exist different functional roles which, depending on the agent and the context, motivating doxastic reasons can play.

In this case, the weatherman's testimony plays the role of what I will call a *direct doxastic reason*. If R is a reason that motivates you directly, then R motivates, sways, or guides you toward a doxastic attitude – which could be belief, disbelief, suspension of judgment, or an increase/decrease in credal state – in a way that is unmediated by other considerations. It sways and would sway you toward a doxastic attitude even if it is the only consideration you consciously entertain.

In contrast, the fact that the weatherman is a casual prankster plays the role of what I will be calling an *indirect doxastic reason*. The role is 'indirect' because reasons which play it in a given episode motivate us indirectly, mediated by other explicit deliberative premises. A purely indirect reason does not function as though it directly favors an answer to the question of whether p is true – it would be incapable of swaying us toward or away from belief considered on its own – but it acquires mediate relevance to deliberation by bearing on the question of how to weigh other reasons. An indirect reason may lead us to increase or decrease the weight of another consideration, or enable or disable (turn 'on' or turn 'off')[7] another consideration from having weight in deliberation entirely.

Of course, some will deny that indirect reasons like the weatherman's prank-pulling

[6]Dancy also discusses enablers/disablers, which – like attenuators/intensifiers – bear on the question of how or whether to weigh other considerations, but which function by turning other considerations on or off. While attenuators/intensifiers attenuate or intensify the weight other considerations carry, enablers/disablers either enable other considerations to carry weight when they wouldn't have otherwise, or prevent ("disable") other considerations from carrying weight when when they would have otherwise.

[7]An example of a disabler: Jones confesses to the crime, which leads Detective Smith to lean toward belief that Jones is guilty. When Smith finds out that the confession was coerced, it becomes essentially worthless as evidence; the coercion disables the confession from carrying any weight in Smith's deliberation at all.

qualify as genuine doxastic reasons. I'll consider objections along these lines later. But I take it that the contrary will be fairly intuitive to most readers, so we'll assume it for now and see what follows.

Now, consider the following cases:

> RACE&IQ: Janet is deliberating about whether to believe that white people are on average more intelligent than black people. Growing up, Janet was taught that IQ scores are accurate, unbiased measures of intelligence. She thinks to herself...
>
> "From what I can tell, the empirical literature suggests that white people tend to score slightly higher on IQ tests than black people. This seems pretty significant. So I guess I'm leaning toward belief. But wait – that's odd – what is this ugly feeling in my stomach? Well...I suppose I've heard Jennifer and Tara say that it's racist to place much significance on the IQ statistics. I can't remember why they said that, but they're pretty much always right when it comes to moral topics. So, they're probably right here: it would be morally wrong to put much weight on the IQ statistics. For this reason, I'll put less stock in that than I was initially inclined to. And now that I'm weighing the data only as heavily as it deserves, it's clear that my total evidence really doesn't merit belief. I'll suspend judgment."

In this case, Janet consciously takes into account both the statistics about IQ and testimony from apparent moral authorities. While the former functions as a direct reason by swaying Janet toward belief independently, the latter does not. After all, the testimony bears just on the question of how to weigh or evaluate the statistics, and so could not have swayed Janet toward any doxastic attitude on its own. It functions instead as an indirect reason, like the weatherman's prank-pulling: it simply guides her to attenuate, or put less stock in, the statistics. This ends up making the difference between belief and suspension of judgment.

It may initially be tempting to imagine that Janet sees the immorality of putting much weight on the statistics as evidence that they are somehow unreliable. However, this interpretation is unnecessary. In fact, let's stipulate that Janet does not see things this way. While Janet is confident that, morally speaking, she shouldn't put much weight on the statistics, she has no idea what to think about how (or if) this connects to their evidential value. All of this might mean Janet was irrational to trust her friends. But the point here is not whether Janet was rational, but whether it is possible for an ordinary person to reason like her. As I have described the case, her trust in her friends plays an important (even if indirect) role in an authentic episode of doxastic deliberation.

Consider also the following case:

> #METOO: Jose, who has been very influenced by the MeToo movement, is disturbed when his good friend Leslie accuses her ex-boyfriend, Fred, who is also a friend of Jose's, of sexually assaulting her. First, Jose speaks to Leslie, and is very moved by her testimony. When he talks to Fred about the issue, Fred denies everything and claims Leslie is trying to get back at

him for breaking up with her. Later, Jose deliberates about what to believe. He reasons roughly as follows:

> "In my gut, I feel that Leslie is telling the truth. She seems more credible, and her testimony feels much more compelling. In contrast, I can't shake the feeling that Fred is lying or faking outrage. So I'm inclined to weigh Leslie's testimony much more heavily than Fred's, enough to suffice for belief. But then again...my emotional reaction to her testimony probably biased my interpretation of Fred's behavior. In fact, if I'd spoken to Fred first, his testimony probably would've seemed more credible, and Leslie's fake! I know if I gave any weight to this, I wouldn't trust how things seem with my evidence and would just suspend judgment. But something about that would feel wrong, and I think I know why: *morally, we ought to believe women about such matters*. For this reason, I think it's actually OK to embrace my gut intuition even though it's biased by non-evidential factors. I believe Leslie."

This case is more complicated. Jose consciously takes into account both Leslie's and Fred's testimony, both of which function as direct reasons. However, Jose's emotions make Leslie's testimony (seem like, function like) a stronger direct reason than Fred's, because it makes her testimony seem to have more evidential value. When Jose takes into account that the way things seem with his evidence was probably influenced by factors unrelated to the truth of the matter, this fact threatens to disable his intuitions[8] entirely, and so functions as an indirect reason.[9] This wouldn't make his intuitions disappear (often, intuitions linger long after we've deemed them unreliable), but it would lead Jose to stop trusting them, and thus to deny them any weight in his deliberation.

Before this happens, Jose considers an apparent moral obligation to believe women, which he essentially takes as permission to lean into his gut feeling despite the fact that it was likely influenced by non-evidential factors. In other words, a moral consideration disabled or "turned off" a potential evidential disabler (e.g., the fact that his gut was biased). This moral consideration functioned as an indirect doxastic reason, in the sense that it only influenced Jose's deliberation by influencing how or whether he weighed another reason.

Like Janet, it may be that Jose's reasoning (or belief-set) is irrational. But in this paper, we're focusing solely on the question of whether it is psychologically possible[10]

[8] Or seemings - it's not important for my argument which of these it is.

[9] If the intuition itself (or the fact that he had the intuition) functions as an indirect reason, then the fact that his intuition was biased functions as a kind of 'second-order' indirect reason – but importantly, still an indirect reason. For my purposes, any reason that functions by modifying how another reason is weighed (whether the reason it modifies is direct or indirect) will count as an indirect reason.

[10] Of course, there may be different senses of 'psychologically possible.' X counts as 'psychologically possible' in the sense I'm interested in here if there are no fundamental constraints on human psychology which make X impossible for all human beings (across all times and places) absent unnatural or extreme interventions – wonder drugs, science-fiction technology, magic, and so on. In this sense, it is psychologically possible to win the Pulitzer Prize in literature, outsmart a grand master in a game of chess, reason according to logical fallacies, and to believe the CIA is out to get you. Psychologically impossible feats include becoming fluent in an entire language from just one two minute conversation, and clearly holding (and understanding) every step in Fermat's Last Theorem in your head at once.

For a more detailed discussion of psychological possibility, see Mason (2003).

to reason as they do. This is not to suggest that psychological possibility and rationality are completely separate. To an extent, rationality may constrain what is possible. For instance, it may not be possible to infer that x is round from the fact that x is square (and so on).[11] Our own sense of rationality may also constrain what can be a motivating reason for us individually. But clearly, if there are rational constraints on psychological possibility, they do not prevent false beliefs, sloppy reasoning, reasoning according to logical fallacies, or mistaking bad or non-reasons for good reasons. So when I say that we are setting aside questions about rationality, what I mean is that we are ignoring the question of whether Jose and Janet have violated thick or substantive rational norms to focus instead on whether it is psychologically possible (according to whatever thin notion of rationality may constrain psychological possibility) for them to reason as described.

Now, we can return to the main argument of the paper:

1. Evidential considerations that play the role of 'indirect reasons' in doxastic deliberation are motivating doxastic reasons.
2. Non-evidential considerations can play the role of 'indirect reasons' in doxastic deliberation.
3. If (1) and (2), then non-evidential considerations can be motivating doxastic reasons.
4. If non-evidential considerations can be motivating doxastic reasons, then exclusivity and transparency are false.
5. (from 1, 2, 3, 4) Exclusivity and transparency are false.

Some of these premises are worth objecting to and others are not. Premise (5) follows from other premises and (4) should be uncontroversial – it follows from the definitions of exclusivity and transparency. We won't explore objections to those. This leaves us with (1), (2), and (3). Premise (1) is highly plausible on its face, as illustrated by WEATHER PRANKS. Premise (3) comes from the thought that functioning like a motivating doxastic reason suffices for being one. Since motivating reasons are psychological rather than normative,[12] I take it that (3) is also prima facie plausible. I'll consider objections to both later, but I assume (2) will be the most contentious. As a result, most of the ink spent on objections will be dedicated to defending it. Premise (2) will be true just in case RACE&IQ and #METOO (and/or cases like them) are psychologically possible – that is to say, possible for some human being(s) without science-fiction pills, magic, or so on. In both cases, agents explicitly consider non-evidential considerations while reasoning about what to believe. These non-evidential considerations influence how their deliberation unfolds and guide them to conclude differently than they would have otherwise, in much the same way that the weatherman's pranks guide Leticia. In other words, both cases involve non-evidential considerations that function as explicit

[11] Partly for this reason, we might say there is a sense in which all agents possess some basic level of rationality. But all plausible theories of psychological state attribution which build this in must be compatible with the fact that there is a thicker sense of the term 'rational' in which irrationality is not just possible but actual.

[12] Of course, this leaves open that our subjective senses of normativity impact what can and cannot be a motivating reason for us individually.

premises in deliberation about what to believe, and so function as indirect doxastic reasons. If these cases are possible, and indirect reasons are motivating reasons, then exclusivity and transparency are false.

I take it that the cases I've presented seem psychologically possible. This intuition of possibility calls out for explanation. In many ways, the most natural explanation – and the one I am offering – is that the cases seem possible because they are possible. Notice the virtues of this explanation. It is simple, unsurprising, and respects the fact that intuitions are generally reliable sources of knowledge. It requires no strange assumptions, skeptical hypotheses, or further explanation. All we must do is assume our mental faculties are functioning normally.

Of course, my opponent will reject this explanation. But before I consider and respond to her objections, I too have an explanatory burden to discharge. In the literature, many intelligent people have assumed transparency and exclusivity are obviously true. If they're false (as I'm suggesting), then why have they seemed obvious to so many people, including those who reject the views they're typically taken to support?[13] For most, I take it that exclusivity (the claim that only evidential considerations can be motivating doxastic reasons) seemed obvious because transparency seemed both obvious and to entail exclusivity. Once again, transparency is the claim that from the first-person perspective in doxastic deliberation, the question of whether to believe that p collapses into – and for our purposes, really just becomes – the question of whether p is true. Because transparency is widely regarded as an uncontroversial feature of deliberation, most authors who reference it do not bother to motivate it. But the few that do (most notably Shah) tend to appeal to the lack of inferential steps between answering the question of whether p is true and answering the question of whether to believe it. For instance, Shah writes:

> It is not as though, in deliberating about whether to believe that p, the reason why one focuses on whether p is the case is that one has noticed that as a matter of psychological fact one has come to believe only what one has ascertained to be the case. This would involve an inferential step: 'Should I believe that p? Well, I shall end up believing that p if and only if I ascertain that p is true, so I had better consider whether p is true.' But there is no such inferential step involved in moving from the question whether to believe that p to the question whether p is true. When I ask myself whether to believe that it is raining, the question whether it is raining becomes immediately and solely relevant. I recognize immediately that the only way to answer the former question is to answer the latter. (Shah, 2006)

In other words, an answer to the question of whether p is true *just is* (at least from the first-person perspective) an answer to the question of whether to believe it, and vice versa. To have answered one automatically means you have answered the other, with no need to infer either from the other.

In one respect this observation is mundane. To my knowledge, no one has challenged the idea that to believe something is at least in part to think it is true. That this

[13] For instance, Foley was a committed pragmatist, but still took doxastic deliberation to be transparent (Foley, 1993).

is clear from the outset of doxastic deliberation isn't terribly interesting – all it seems to entail is that we know going in that the conclusion of our deliberation will be an attitude toward truth. This is compatible with arriving at such an attitude at least in part by considering factors we see as unconnected to p's truth, but which bear on other things that remain relevant either to the question of what to believe or to the question of how to reason. So on its own, this doesn't get us transparency.

However, it does help pave the way for another kind of intuition-pumping which appeals to the idea that the only way to make headway or progress in doxastic deliberation is by trying to answer the question of whether p is true.

> If my answering a question is going to count as deliberating whether to believe that p, then I must intend to arrive at belief as to whether p just by answering that question. I can arrive at the belief just by answering the question whether p; however, I cannot arrive at the belief just by answering the question whether it is in my interest to hold it. (Shah, 2006)

In this passage, Shah seems to be pointing out that the only way to make progress in doxastic deliberation is by considering factors we take to bear on the question of whether p is true. To make progress in deliberation is to be swayed by a reason (or reasons) toward or away from a particular conclusion. Notably, we can make progress in deliberation without thereby reaching a conclusion; it is sufficient to be swayed. For example, very weak evidence might sway me slightly toward belief even if in the end I deem it insufficient and suspend judgment. So Shah's reasoning here seems to be something like this: I can make progress in deliberation about whether to believe p *just* by considering factors that bear on p's truth. In contrast, I can't make progress *just* by considering factors I see as unrelated to p's truth, like my desires or interests. So considering evidence is both necessary and sufficient for making progress in doxastic deliberation, while considering non-evidential factors is neither necessary nor sufficient.

This still doesn't get us either transparency or exclusivity. It's just an appeal to our intuitions about what can and cannot be a direct doxastic reason. Recall that direct reasons immediately engage the question of whether p is true, and so sway us toward believing p (e.g., help us make headway in answering the question of what to believe) all on their own, in the absence of other conscious considerations. Shah seems to have ignored the existence of indirect doxastic reasons.

Here's a competing suggestion: when the evidential situation is uncomplicated – when only direct reasons come to mind and we don't need to think twice about how to evaluate them – the question of whether to believe p really is transparent to the question of whether p is true. In such cases, treating a consideration we see as unconnected to p's truth as a direct reason for belief would require us to simultaneously treat it as immediately relevant to the question at hand (that of p's truth), while also acknowledging it as completely irrelevant (e.g., as non-evidential) to the very same question. Thus, to regard a non-evidential consideration as a direct reason for belief, particularly in uncomplicated evidential situations, may seem to require a level of delusion.

But what about indirect reasons? It's worth noting that we cannot make headway in answering the question of whether to believe it will rain *just* by considering those,

either. The question they bear on (how to weigh other reasons) is engaged only after direct reasons are in play. So Leticia could not make progress in answering the question of whether to believe it will rain purely by considering the weatherman's prank-pulling. The pranks become relevant only after she considers his testimony. This doesn't mean the prank-pulling isn't an explicit premise in Leticia's reasoning; it just means it's not in that instance a direct reason.

Something similar could be said of non-evidential considerations. Because non-evidential considerations cannot bear on the question of whether p is true by definition, perhaps they cannot be direct doxastic reasons. But this does not mean they have no role to play. As cases like #METOO and RACE&IQ illustrate, non-evidential considerations can still (from the first-person perspective) engage the question of how/whether to weigh other doxastic reasons, and guide deliberation accordingly, once other considerations have been brought to bear. That they cannot do this on their own does not mean they cannot (in some contexts) do so by playing indirect roles. After all, when indirect reasons are in play, we are considering the question of how to weigh other reasons, which is an entirely separate question from that of whether p is true. While it may be tempting to think that the question of how to weigh other reasons is somehow subservient to the truth-question, examples like #METOO and RACE&IQ indicate otherwise.

In other words, the intuition-pumping which led so many to think that transparency and exclusivity were obvious focused solely on oversimplified cases of doxastic deliberation that involved only direct reasons – those considerations which, on their own, can help us make progress in answering the question of what to believe. While this all lends support to the idea that only evidential considerations can be direct doxastic reasons, it's irrelevant to the question of what can be an indirect doxastic reason. This is an easy mistake to make in a literature that does not emphasize the distinction between direct and indirect reasons. But once we attend to it, and to cases like #METOO and RACE&IQ, we can see that transparency has little ground to stand on.

4 Objections

In #METOO and RACE&IQ, non-evidential considerations feature as premises in Jose's and Janet's doxastic deliberation. While they do not guide either agent toward or away from a particular conclusion in isolation, they function as indirect doxastic reasons because Jose and Janet take them to bear on the question of how to weigh other reasons. While both #METOO and RACE&IQ may be rationally defective, they appear on their face to be psychologically possible. I've offered an explanation of this intuition: they seem possible because they are possible. My opponent will reject this explanation, but she needs to provide grounds for doing so. Correspondingly, I anticipate three main camps of objectors. The first deny that the cases I've presented pose any threat to exclusivity, and instead reject premise 1 or premise 3. I discuss these in 4.1. In 4.2, I explore objections that target premise 2.

4.1 No Threat Objections

No-Threat objections accept that #METOO and RACE&IQ as described are psychologically possible, and so accept premise 2. However, they deny that this poses a genuine threat to exclusivity or transparency. This could be because indirect evidential reasons are not genuine doxastic reasons, or because the functional similarities between evidential and non-evidential indirect reasons do not guarantee that both are motivating reasons. Correspondingly, there are two types of No-Threat objection.

4.1.1 Denying Premise 1

The first denies premise 1, which claims that indirect evidential reasons can be motivating doxastic reasons. On a first pass, some might worry that we don't need the distinction between direct and indirect reasons because we already have the language of defeaters. But this doesn't make sense upon reflection. If the point is supposed to be that a fact cannot be both a defeater and a doxastic reason for the same proposition, then it's false; that storm clouds are gathering overhead is both a rebutting defeater for and a direct reason to disbelieve the proposition that it will rain today. Presumably something similar could be true of indirect reasons and undermining defeaters. In fact, that the distinction between rebutting and undermining defeaters is well-recognized in other literature is excellent evidence that it's a meaningful distinction worth capturing in our terminology. The language of defeaters typically crops up in discussions about justification, where defeaters are understood as either propositions the subject ought to believe (but may or may not believe in actuality),[14] or mental states[15] that interfere with subjects' abilities to have knowledge. In contrast, motivating reasons tend to come up in discussions about action and reasoning, and are most often construed as belief-desire pairs,[16], facts we believe,[17] or propositions we believe.[18][19] These categories are not mutually exclusive. Whether it makes sense to speak in terms of defeaters or reasons depends on the literature we're situated in and the questions we're asking (e.g., whether subjects meet the conditions for knowledge and blame, versus the psychology of reasoning and agency). Since the focus of this paper is the psychology of reasoning and agency, I'll continue to speak of reasons.

A more compelling version of the No-Threat objection might include people like Jonathan Dancy,[20] who thinks that indirect reasons do not actually favor an answer to

[14]For examples of this view, see Harman (1980), Pollock (1986), Meeker (2014), Lackey (2014, 2016), Goldberg (2017, 2018), and Benton (2016).

[15]For an example of the mental state view of defeaters, see Bergmann (2005, 2006).

[16]Examples of *psychologism* in the literature include (Davidson (1963), Smith (1994), Wallace (2003), Pryor (2007), Turri (2009), Gibbons (2010)

[17]Examples of this view, known as *factivism*, in the literature include Hyman (1999, 2010, 2015), Bittner (2001), Stout (2009), Alvarez (2010, 2018), Raz (2011a), Littlejohn (2012), and Roessler (2014).

[18]Examples of *non-factivism* in the literature include Parfit (1997); Dancy (2000, 2003, 2011); Miller (2008); Schroeder (2008); Setiya (2011); Marcus (2012); Comesaña and McGrath (2014); and Singh (2019). Though Dancy notably prefers to speak in terms of 'prospects' (facts or apparent facts) rather than propositions, his view is still non-factive, and prospects will count as propositions according to many theories of propositions.

[19]For a recent discussion of these distinctions, see Cunningham (2021).

[20]Much thanks to Jonathan Dancy for his helpful comments on this section.

the question of whether to Φ, and so are not genuine reasons at all. Dancy terms what I have been calling 'indirect reasons' 'attenuators/ intensifiers' and 'enablers/ disablers' instead (Dancy 2000, 2018).

An adequate response to this objection could be a paper in its own right. For now, what's important is two things. First, even if this saves exclusivity, my argument has still proven that transparency is false (more on this below). Second, what's most important for my argument is simply that it's plausible indirect reasons are genuine doxastic reasons. The topic of what exactly it means to be a motivating reason is controversial,[21] but what should be uncontroversial is that direct reasons are paradigmatic examples. So while I cannot hope to provide a precise definition of motivating reasons, I can highlight important features that direct and indirect reasons have in common to motivate the idea that they're worth grouping together.

First, we clearly see indirect reasons as having normative significance to the question we are deliberating on. That is to say, we are conscious of indirect reasons, like the weatherman's history of pulling pranks, as things that call out to be weighed and contended with in reasoning about whether to believe that p, as opposed to reasoning about whether to hope that p or believe that q. That we are conscious of them in reasoning makes them quite different both from factors that only influence deliberation subconsciously (ex: biases, inference rules) and from factors we may be conscious of but see as irrelevant to the question of whether to believe that p (ex: background noise).

Second, direct and indirect reasons can both make a difference to how doxastic deliberation on p actually concludes. Just as the weatherman's testimony can lead me to believe it will rain, his history of pulling pranks can lead me to conclude deliberation in a doxastic attitude I would not have concluded in otherwise, like suspension of judgment, or a subjective credence of 0.6 rather than 0.8. While his prank-pulling could not make a difference on its own, it makes a distinctive contribution when paired with the right direct reasons.

Third, both types of reason guide us to move closer to a doxastic attitude even when they do not make a difference to how our deliberation actually concludes. We may think of this as adding to or subtracting from the total weight or balance of the reasons without making the difference between one doxastic attitude or another. For example, your neighbor's testimony that rain is on the way may lead you to move slightly toward belief that it will rain, more than you would have leaned otherwise, even if it does not make a difference to the actual outcome of your deliberation because you had stronger reasons to think the opposite. Important to this third point, too, is that direct and indirect reasons both guide us toward or away from doxastic attitudes (belief, disbelief, suspension of judgment, credal states) with respect to p, as opposed to other types of attitudes (ex: hope that p, outrage that p) or doxastic attitudes on other topics (ex: belief that q).

Finally, the second and third feature noted above are both true by virtue of the first. In other words, the weatherman's pranks guide Leticia away from concluding in belief *because* she sees them as something to be weighed and contended with in reasoning about whether to believe it will rain, i.e., *because* she takes the pranks to have normative significance for the question she is deliberating on. Presumably, it is also by virtue of

[21] See Cunningham 2021 for a recent discussion.

this that the testimony manages to guide her toward belief.

While this doesn't show conclusively that the kinds of considerations I've been calling 'indirect reasons' are genuine instances of motivating reasons for belief, it does lend plausibility to the idea that they are. Those who remain unconvinced can take this paper as conditional on that assumption going forward, and will find that it still has interesting implications. First, if exclusivity only applies to direct reasons, then there is a large and interesting class of considerations exclusivity does not apply to, but which share many or most interesting features with motivating reasons for belief. This makes exclusivity less interesting than we've typically supposed, and makes the possibility of cases like #METOO and RACE&IQ significant even if 'indirect reasons' are not technically motivating doxastic reasons. Second, and more importantly, it still entails that transparency is false. Jose and Janet explicitly regard considerations they see as unconnected to p's truth as relevant to the question of whether to believe that p during doxastic deliberation. If transparency is right, this should be impossible: the only considerations we can see as relevant should be those we see as bearing on the question of p's truth.

4.1.2 Denying Premise 3

The second type of No-Threat objection targets premise 3. Objectors in this camp accept that indirect evidential reasons can be motivating doxastic reasons (premise 1), and also accept that #METOO and RACE&IQ are psychologically possible (premise 2, or the support for it). Instead, they take issue with the idea that *functioning* like a motivating doxastic reason suffices for *being* one. Objectors in this camp will appeal to the idea that motivating doxastic reasons are not merely things that function as premises in deliberation – they are also things that are true, or that justify our beliefs.

The idea that motivating reasons are factive is common, and of course, many will think the non-evidential considerations in #METOO and RACE&IQ are not facts. If they aren't, then perhaps they can't be motivating reasons, either. Rather than taking a stand on this issue, I'll simply point out that non-factual considerations can still be apparent motivating reasons, or considerations that would count as motivating reasons if only they were true. If this paper succeeds in proving that non-evidential considerations can be apparent motivating doxastic reasons, then I'll consider it a win.

The other version of this objection stipulates that motivating doxastic reasons are things we can cite to justify our beliefs, in addition to things that can function as explicit premises in our reasoning. Clearly, there is some way to interpret this requirement which would make non-evidential considerations incapable of being motivating doxastic reasons: if what is meant by 'cite' is 'legitimately cite,' if the notion of 'justification' at issue is purely epistemic, if epistemic justification can only come from evidential factors, and if motivating doxastic reasons must be able to justify belief in this sense specifically, then non-evidential considerations cannot be motivating doxastic reasons. Of course, the issue with this response is that it's blatantly question-begging.

Broadly construed, there are many ways to interpret the requirement that motivating doxastic reasons be capable of justifying the doxastic attitudes they support which don't beg the question. None of these interpretations threaten my argument. For one, there are many non-epistemic notions of justification; plausibly, non-evidential con-

siderations could justify doxastic attitudes non-epistemically (ex: morally). It's also not a settled matter that epistemic justification can only come from evidential factors. The growing body of literature on moral and pragmatic encroachment includes many authors who take seriously the idea that moral or pragmatic factors can influence the thresholds required for[22] or even contribute to[23] epistemic justification, in which case non-evidential considerations may contribute to epistemic justification.

Alternatively, if the complaint is that no one would ever cite exclusively non-evidential factors to justify their belief if challenged, then there are two things worth noting. First, what we cite when challenged to justify our beliefs is heavily context-dependent. If Leticia's friend asked her why she's suspending judgment about rain after seeing the forecast together, Leticia probably wouldn't mention the weatherman's testimony even though it is among the reasons for which she is suspending judgment (what's the point, if her friend already knows about it?). Instead, she's more likely to cite the weatherman's history of pulling pranks. In other contexts, she might cite the testimony itself while leaving out the pranks. In others in might be her location, the local climate, the fact that the weatherman has expert knowledge, the frequency at which he pulls pranks, etc. In general, the considerations we cite when challenged to justify our beliefs depend heavily on context, and especially on the body of assumptions we take ourselves to share with those present.

Second, it simply isn't true that no one would ever cite non-evidential considerations when challenged to justify their beliefs. If a Jordan Peterson fan challenged Janet to explain why she is suspending judgment on the question of whether white people are more intelligent than black people in spite of the statistics about IQ, it would be perfectly natural for Janet to reply, "I think it's immoral to put much weight on those statistics." Similarly, if someone aware of the evidential situation asked Jose why he believes that Fred sexually assaulted Leslie, it would be natural for Jose to reply, "I think it's morally important to believe women on such matters." This was a common response among supporters of the #MeToo movement when challenged to justify their belief that Supreme Court Justice Brett Kavanaugh sexually assaulted Dr. Christine Blasey Ford.[24] So there are plenty of contexts where we'd expect ordinary people to cite non-evidential considerations when pushed to justify their doxastic attitudes.

4.2 Debunking Objections

In the previous section, we considered objections to premise 1 and premise 3. Starting in this section, we will turn our attention to premise 2, which I assume is the more controversial premise. Recall that premise 2, which claims that non-evidential considerations can function just like indirect evidential reasons in doxastic deliberation, will be true just in case scenarios like RACE&IQ and #METOO are psychologically possible. In both cases, agents who are deliberating about what to believe explicitly take into account non-evidential considerations that act as premises in doxastic deliberation. If these cases are possible (and premises 1 and 3 are true), then transparency is false.

[22]For instance, Basu & Schroeder (2019), Gardiner (2018), and Moss, (2018).

[23]For an interesting defense of this claim, see Gao, (2019).

[24]For an overview, see Cauterucci (2018).

Again, I take it that #METOO and RACE&IQ seem psychologically possible at least on their face. While this intuition could be illusory or confused, it's a datum that needs to be explained. The explanation I am offering is that the cases seem possible because they are possible. This explanation has many virtues. It is simple, natural, and unsurprising. In many ways, this is the default explanation – to have an intuition that p when p is false usually means something unusual has happened, which itself calls out for explanation. Because intuition is a generally reliable source of knowledge, my explanation of our intuitions about #METOO and RACE&IQ does not raise more questions than it answers. All we must do is assume that our mental faculties are working normally.

On the face, the most tempting objections are simply attempts to debunk these intuitions, or my explanation of them. A successful debunking objection does two things. First, (1), it offers a plausible alternative explanation of our intuitions that does not threaten transparency – it explains why the cases I've presented seem possible if in reality they are not. But as I've noted, my explanation has many virtues, and so could still be the best explanation even if alternative explanations exist. So in addition, the debunker must (2) provide reason for thinking that her explanation is equal to or better than mine. As it turns out, the first criterion is not difficult to satisfy, but my opponent will run into trouble in the second.

The debunker I'm imagining believes that our intuitions have been confused or muddled by the possibility of similar cases that are compatible with transparency. One thought is that #METOO and RACE&IQ only seem psychologically possible because they resemble cases in which agents are deliberating about what to *accept*, rather than about what to *believe* For example, someone in a situation very like Janet's may think she's morally obligated to live her life as though she is not confident that white people are more intelligent than black people, even if secretly she is confident of that. Or similarly, Jose may think he's morally obligated to accept Leslie's testimony for the purposes of action even if he is uncertain about its truth.

The problem with this suggestion is that it doesn't have much prima facie plausibility. Paradigmatically, acceptance is an attitude that is action-oriented rather than truth-oriented. For that reason, we'd expect deliberation about what to accept to involve, at least in part, deliberation about what action(s) to take in some context(s). But Jose's and Janet's deliberations include no mention of specific actions. Nor do they mention anything we'd expect action-oriented deliberation to be concerned with on such topics, like the moral risks of acting one way rather than another. Instead, both agents' deliberations are concerned exclusively with truth, evidence, and moral obligations pertaining to belief and the assessment of evidence (rather than moral obligations pertaining to action). This isn't to say that there couldn't be some form of deliberation about what to accept that resembles #METOO or RACE&IQ more closely. But it would be a stretch. The burden of finding such a case (and then explaining why we shouldn't prefer my explanation over the debunking one) is on my opponent.

A more plausible debunking explanation suggests the intuition that #METOO and RACE&IQ are psychologically possible actually arises from similar, transparency-friendly cases where evidence is the sole focus of agents' attention. After all, it's plausible in both cases that the evidence objectively favors the doxastic attitudes Jose and Janet form. For this reason, it is easy to imagine cases very similar to Jose's and

Janet's in which agents come to the same judgments based on just their evidence. For example, someone in Janet's situation may have taken her friends' testimonies to be evidence that IQ tests are unreliable. An individual in Jose's position may have thought the perceived moral imperative to believe women is evidence of the fact that it's uncommon to lie about sexual assault. While these situations are incompatible with #METOO and RACE&IQ as described, they resemble the cases in important ways.

This explanation also pairs nicely with the uncontroversial idea that non-evidential factors, like moral beliefs, desires, and biases, can influence doxastic deliberation subconsciously. This could easily happen in situations like Jose's and Janet's. The belief that morally we ought to believe women could subconsciously influence someone like Jose to reason sloppily, or to dismiss the realization that his gut was biased by non-evidential factors too quickly. But nonetheless, my opponent would maintain that #METOO and RACE&IQ *as described* are psychologically impossible. While we can be influenced by non-evidential factors subconsciously, we cannot consciously acknowledge them as relevant to our deliberation.

This is an alternative to the explanation I've offered, so the debunker has satisfied criterion (1). The more difficult task is satisfying (2): explaining why my explanation is not superior. At a first glance, my opponent may claim that the mere existence of similar, transparency-friendly cases is a reason in itself against preferring my explanation over theirs. The thought is that the resemblance #METOO and RACE&IQ bear to cases that aren't problematic for transparency makes our intuitions about #METOO and RACE&IQ so unreliable that our intuitions hardly count as evidence against transparency at all.

However, on reflection, this is way too strong. Appeal to cases is an important argumentative strategy in philosophy, and most cases resemble others in various respects. That a problematic case has a number of things in common with unproblematic ones isn't typically a reason to think our intuitions about the former are totally unreliable. If it were, far less ink would be spilled on ailing violinists, philosophical zombies, twin earths, and real barn sightings in fake barn country – which could all be "debunked" in much the same manner my opponent is attempting to debunk #METOO and RACE&IQ. To that end, my opponent's debunking explanation lacks explanatory power; it fails to apply in a wide range of relevantly similar contexts, including almost every example and thought experiment that has ever posed a problem for a view in philosophy.

My opponent must do better: she needs to find a debunking explanation that doesn't also apply to swampmen or sick violinists. One place she may turn is our current social and political climate. A hundred years ago, long before the #MeToo movement and heated debates about race and IQ, my examples might have seemed much less plausible, or even silly. In contrast, there seems to be something timeless about our intuitions regarding zombies, twin earths, and so on. Maybe this is evidence that our intuitions about the cases I've presented are being affected by factors unrelated to the cases' possibility.

While I acknowledge that intuitions about #METOO and RACE&IQ may have been different a hundred years ago, I don't think this is evidence that our current intuitions are unreliable. A hundred years ago, both cases probably were psychologically impossible for most people given their normative beliefs. Consider again the fact that

a consideration cannot be a motivating doxastic reason in the relevant sense unless one sees it as a normative doxastic reason. A hundred years ago, very few people would have recognized the existence of exclusively moral normative reasons for believing women about sexual assault or for not putting much stock in IQ statistics. Correspondingly, such considerations could not possibly have featured as premises in most people's doxastic deliberation. In other words, our current cultural climate may be causally responsible for our intuitions about #METOO and RACE&IQ, but only because it has changed many of our normative beliefs, which in turn has changed what is psychologically possible for most (or many) people.

At this point, I think it is clear that debunking-style explanations of our intuitions about #METOO and RACE&IQ face a potentially insurmountable task. While it is easy to propose some alternative explanation, it is difficult to explain without begging the question why the simple explanation I have offered is not the best available. I leave it open that this is possible, but clearly, it would require significant argumentative work.

Before we wrap this section up, there is one way to make debunking objections more plausible – and that's to appeal to transparency itself. While this looks a bit question-begging, it is widely assumed to be true in the literature. If it were established fact, that would make debunking explanations explanatorily superior to the explanation I've offered.

So let's take a moment to reflect on the support for transparency and exclusivity. Because both claims concern our psychology, neither can be derived from logical or mathematical axioms. They cannot conflict with or be supported by established laws of physics, chemistry, or biology. Theoretically, one or both could be supported by psychological research. But any study that could be conducted to verify them would have to rely heavily on subjective reports about deliberation, which philosophers are unlikely to trust. And either way, empirical research has not been cited specifically in support of either claim in the philosophical literature. From what I can tell, that's because no such research exists. What's worse is that neither claim has been argued for.[25] Both are typically assumed in the literature without either argument or empirical support. The reason for this, presumably, is that most of those writing on the topic have thought one or both claims are self-evident. At a first glance, we might think self-evidence is a matter of obviousness. But of course, the whole point of the cases I've presented is that transparency and exclusivity are not obvious. Alternatively, we might think of self-evidence as obviousness to anyone who fully understands the meaning of the terms in question.[26] If doxastic deliberation is by definition deliberation that is transparent only to the question of whether p is true, then transparency is self-evident, and exclusivity follows. The problem with this move is that the cases I've presented cast just as much doubt on this definition of doxastic deliberation as they do on exclusivity and transparency.

So as it stands, exclusivity and transparency are not self-evident. They are not supported by existing scientific literature and have not been argued for. They are unverified empirical conjectures. There is no reason to prefer a favor a debunking explanation simply because it would be friendlier to them.

[25] Psychological exclusivity has been argued for, but only from transparency (Shah, 2006).

[26] See Audi 2019 for a recent discussion of this view of self-evidence.

4.3 Absurd Implications

Another class of objections to premise 2 turns on the concern that these cases generate absurd implications. It may seem that if the examples are possible, then Jose and Janet will have strange answers to questions about p's objective probability. Perhaps Janet will end up having to say that, "The probability that white people are on average more intelligent than black people is high, but I think it is not high." In general, it seems absurd to imagine agents who say things like:

> "The objective probability that p is really high, but I think it is low."
>
> "The evidence strongly supports p, but I think it doesn't."
>
> "I think p's probability is significantly higher than its objective probability."
>
> "I think p is highly probable, but I don't think it's reasonable to think this."

I agree that these statements are troubling; it's very difficult to imagine any agent assenting to them honestly. So it would be a problem for my examples if their possibility entailed the possibility of agents making such absurd reports in good faith.The most troubling feature of these reports is that they seem to reflect two settled, contradictory attitudes related to p's probability: one which takes into account only the evidence, and another that factors in non-evidential considerations. These attitudes could be credences, beliefs about p's objective probability, beliefs about evidential support, or higher-order beliefs about what is reasonable to believe.

On the one hand, if we must interpret my examples as involving conflict between agents' credences and higher-order beliefs about what is reasonable to believe, then the possibility of my examples would entail the possibility of epistemic akrasia, which is controversial in itself. Alternatively, my opponent may think my examples entail that it is possible to have two contradictory credences that p, or that agents could have subjective credences toward p that conflict with their first-order beliefs about p's objective probability. Both of these seem absurd.

With this in mind, #METOO and RACE&IQ will pose problems if we have to interpret Jose and Janet as forming two conflicting attitudes toward either p's probability or the rationality of believing p. That is to say, the examples will be problematic if Jose and Janet arrive at settled conclusions about p's probability based on just the evidence, and separately arrive at conclusions about what doxastic attitudes to take toward p, or about what is reasonable to believe.

Thankfully, we need not assume anything like this. Before Jose and Janet conclude, they have not formed any settled answer to the question of what doxastic attitude to take toward p. They certainly have not settled on answers to questions about p's objective probability, or on answers to questions about what is reasonable to believe. Once they conclude (having taken non-evidential factors into account), they arrive at just one doxastic attitude about p's probability, and one (if any) matching higher-order judgment. Thus, we can expect their reports about p's objective probability or about what is reasonable to believe to match their reports about their subjective credences. Janet will insist the objective probability that white people are on average more intelligent than black people is not high at all – and that she is reasonable to think so – and similarly,

Jose will claim there is a high objective probability Leslie was sexually assaulted, and that he is rational to think so.

But how could this be, if Jose and Janet are taking non-evidential considerations into account? The answer is simple: the non-evidential considerations inform their judgments about how to weigh their evidence, and thus, their assessments of both p's objective probability and the reasonableness of believing p. In other words, Jose sees his moral commitment as giving him permission to put stock in his intuition. Once he does this, he factors his intuition into the weighing calculus, and arrives at the single settled conclusion that p is objectively highly probable given his evidence. Similarly, Janet sees an apparent moral fact as a reason to avoid placing much significance on certain empirical data. She responds to this apparent reason by weighing the data less heavily than she was initially inclined, and only then comes to a settled conclusion about p's objective probability.

Importantly, this does not mean that Jose and Janet are reasoning well. It is compatible with my argument that they are straightforwardly wrong, bad, irrational, and even foolish to think the objective probability of p is anything but what they would have thought had they considered their evidence alone.

Some of my opponents may still have lingering worries. Take Jose, who was said to be leaning toward suspending judgment (or alternatively, leaning toward updating his credence that p from very low to 0.5), and would have done so had he not taken the non-evidential consideration into account. What is this "leaning," if not an attitude about p's objective probability given the evidence? And if it is that, then isn't Jose forming additional, conflicting attitudes when he concludes differently on the basis of non-evidential considerations?

For the sake of argument, grant that "leaning" in these cases is an attitude about p's objective probability. In both examples, agents stopped "leaning" one way and began "leaning" differently after factoring in non-evidential considerations. So if "leaning" is an attitude about p's objective probability, then non-evidential considerations led the agents in my examples (however indirectly) to change their minds about both p's objective probability and the reasonableness of believing p.

Was this rational? Perhaps not. Perhaps the rational thing to do would have been to settle on the doxastic attitudes they were "leaning" toward before they took non-evidential factors into account, so their attitudes about p's objective probability and their attitudes about what is reasonable to believe were formed just by consideration of those things that actually bear on objective probability. But as it is, the topic under discussion is not rationality, but psychological possibility.

Notably, it is also compatible with my argument that these examples are psychologically possible only because of cognitive limitations. After all, the question Jose was deliberating on – the question of whether to believe Fred sexually assaulted Leslie, or if you will, the question of whether this is true – is a few steps removed from the question of how or whether to weigh the realization that his intuition was biased by non-evidential factors. Perhaps this distance from the overarching deliberative question together with personal bias or limitations on time, attention, and working memory enabled Jose to make mistakes he wouldn't have made if he had infinite time and cognitive resources. Again, this is perfectly compatible with my argument, which concerns what is in fact possible for human beings to do – not what would be possible if only we

were less cognitively limited.

But perhaps this raises a different worry: I've said that in both examples, agents were "leaning" one way and began "leaning" another in response to non-evidential considerations. We might wonder if this means they are, in effect, *treating* those "non-evidential" considerations as evidential.

This raises substantive questions about what it means to "treat" a consideration as evidential. But let's assume for the moment that this is right: Jose and Janet treat apparent moral facts as evidential. We can then distinguish between a consideration's *being* evidential and being *treated* as though it is evidential. As we've defined it, a consideration C counts as evidential with respect to p for some agent when she sees or believes it to be connected to p's truth or to the question of whether p is true. By stipulation, Jose and Janet see their moral commitments as unconnected to the question of p's truth, so according to this definition, their moral commitments count as non-evidential.[27] Nonetheless, it could be that both individuals "treat" the relevant considerations as evidential simply because they appear as premises in their deliberation.

This may raise another question: what about the indirect 'evidential' consideration that moves Leticia in WEATHER PRANKS? It functions very similarly to the indirect 'non-evidential' considerations in #METOO and RACE&IQ. And like the non-evidential considerations, there's a sense in which Leticia sees the weatherman's prank-pulling (at least on its own) as unrelated to the question of whether rain is on the way. This could lead to the confusing appearance that 'non-evidential' considerations should be in the same category as the indirect 'evidential' considerations like those in WEATHER PRANKS.

An 'evidential' consideration as we've defined it in this paper is a consideration the agent takes or believes to bear on the question of whether p is true. By definition, no agent can take a purely indirect reason to bear directly on the question of p's truth. But she may take it to bear on this question indirectly, in which case it will count as evidential for her. For instance, while the weatherman's pranks are not independently relevant to the question of whether rain is on the way, they do seem to bear on this question in a way that facts about the job market, geometry, and my brother's taste in cars do not. That is to say, while the weatherman's pranks are not immediately relevant to the truth-question, they still seem to have a kind of mediate relevance in virtue of bearing on the weatherman's credibility.

One way of putting the point is that all indirect reasons (are taken to) bear directly on the question of how to weigh other considerations. But not all indirect reasons (are taken to) bear even indirectly on the question of whether p is true (though presumably they must bear, at least indirectly, on the question of whether to believe p). As far as Jose and Janet are concerned, the apparent moral facts in #METOO and RACE&IQ stand in approximately the same relation to the question of whether p is true as the fact that 2+2= 4: all are irrelevant. Jose and Janet take the apparent moral facts to bear on the question of how to weigh other considerations, but they may do so without taking them to bear on p's truth.

Perhaps my opponent can turn instead to common intuitions that (one may think)

[27] Once again, we're understanding the terms 'evidential' and 'non-evidential' in this way out of charity for my opponent.

are inconsistent with the idea that non-evidential considerations can be reasons for which we believe. Consider the following case:

> FUN WEATHER: Susie is deliberating about whether to believe it will rain later today. She thinks to herself...
>
> "Grandma said she knows it will rain today because her bones are aching. But Grandma's only right about fifty-one percent of the time, so I won't put too much weight on what she says. So I'm leaning toward disbelief...but you know what, it would actually be really fun to believe it will rain today. For that reason, I'll go ahead and put extra stock in Grandma's testimony, and conclude in belief."

Unlike the cases I've presented so far, FUN WEATHER seems psychologically impossible. Clearly, we can't just increase or decrease the weight of evidence for any reason at all. This, too, is a datum that needs to be explained. And transparency seems to offer a tidy explanation: Susie can't just assign more weight to the testimony of her 51% reliable grandmother for fun because only considerations we see as connected to p's truth can be reasons for which we believe. And of course, no agent in their right mind would see the entertainment value of believing it will rain as connected to the truth of whether rain is on the way. If we reject transparency and exclusivity, then aren't we committed to the absurd claim that FUN WEATHER is psychologically possible?

By 'psychologically possible,' what is meant is something like 'possible for human beings without unnatural or extreme interventions.' Clearly, FUN WEATHER is psychologically impossible for anyone with normal or reasonable beliefs, and we can explain why without appealing to transparency. Recall that a normative reason to Φ is a consideration that objectively favors Φ-ing, or that is an objectively good reason to Φ. A motivating reason to Φ is (at minimum) a consideration that can be an explicit premise in reasoning about whether to Φ. A consideration cannot be a motivating reason for you to Φ in the way I'm interested in unless you see it as a normative reason to Φ. That is to say, a consideration cannot function as an explicit premise in your doxastic deliberation unless you see it as an objectively good reason to believe (disbelieve, increase/decrease credence, etc.).

Simply put, the entertainment value in believing it will rain seems like a stupid reason to increase the weight of Grandma's testimony. It's difficult to imagine anyone would think otherwise, which makes it difficult to imagine anyone would see the entertainment value as a normative doxastic reason – which, in turn, makes it difficult to imagine the entertainment value could be an explicit premise in anyone's doxastic deliberation. In general, this is why we cannot increase or decrease the weight of evidence in doxastic deliberation for just any reason at all. Not all considerations strike us as normative doxastic reasons, and so not all considerations can motivate us accordingly. This neatly explains the seeming impossibility of FUN WEATHER, and is consistent with the falsity of exclusivity. It also goes some way in explaining why #METOO and RACE&IQ seem possible in a way that FUN WEATHER does not. Even if we do not share Jose's or Janet's moral beliefs ourselves, it is not difficult in our current political climate to imagine agents who genuinely think there are exclusively moral reasons to

believe women about sexual assault, or to refrain from putting much stock in the statistics about race and IQ – whereas it's difficult, if not impossible, to imagine anyone seriously thinking the entertainment value of a belief is a good reason to believe.[28]

Of course, that still leaves open the question of whether FUN WEATHER is psychologically possible for someone with idiosyncratic normative beliefs. Fortunately, rejecting exclusivity need not commit us to a particular answer. On the one hand, if FUN WEATHER is possible for someone with very strange beliefs, then we still have a straightforward explanation of why it seems otherwise: FUN WEATHER seems impossible because it's hard to imagine anyone actually having the absurd normative beliefs it would take to make the case possible. But we need not commit ourselves to this answer. For it could be that the norms governing doxastic deliberation are sensitive to moral factors, but not prudential factors like entertainment value – in which case, FUN WEATHER would be psychologically impossible no matter what Susie's beliefs are. Alternatively, we could say that non-evidentital factors can modify other reasons in a way that leads us to change our doxastic attitudes only when the evidential situation is complex enough that we can become distracted and make mistakes, or when subconscious biases make non-evidential considerations feel unduly compelling. Any of these things could explain why FUN WEATHER is psychologically impossible – and why #METOO and RACE&IQ are not – even if transparency and exclusivity are false.

5 Summary and Conclusion

Doxastic transparency has struck many as obvious, but it's incompatible with the commonsense intuition that moral considerations may be among the reasons for which some people believe women about sexual assault, refrain from taking statistical evidence about race too seriously, believe loved ones about their gender identity, or give a friend the benefit of the doubt. How do we explain this? Reflection upon the intuitions which motivated transparency reveals these were really intuitions about simple cases of deliberation that feature direct reasons only. While it's possible that non-evidential considerations cannot be direct reasons, this does not mean they cannot be reasons for which we believe in a more limited capacity, by playing the role of indirect reasons.

This undermines a number of arguments for views about the aim/norm of belief which rely on transparency, as well as Shah's argument for normative evidentialism. Shah's argument turned on transparency together with the premise that a consideration cannot be a normative reason for belief unless it can be a motivating reason for belief. Either transparency, this premise, or the claim that transparency entails exclusivity is false if non-evidential considerations can be motivating doxastic reasons. The argument of this paper also suggests that we can weigh a variety of aims against each other in doxastic deliberation (e.g., moral aims against truth aims), which makes reasoning

[28] Ironically, this also means that someone who is committed to evidentialism with all their heart and soul will find that their own deliberation is indeed psychologically exclusive. Because they do not see non-evidential considerations as normative reasons for belief, they cannot be swayed by them in doxastic deliberation. Thus, there is a sense in which a commitment to normative exclusivity can create the illusion of psychological exclusivity for evidentialists who reflect upon what is psychologically possible in their own deliberation.

about what to believe importantly analogous to practical reasoning. In turn, this casts doubt on certain forms of doxastic involuntarism, and suggests a much more expansive picture of the kind of responsibility we can have for our beliefs.

6 References

Alvarez, Maria, & Way, Jonathan. "Reasons for Action: Justification, Motivation, Explanation", *The Stanford Encyclopedia of Philosophy* (Fall 2024 Edition), Edward N. Zalta & Uri Nodelman (eds.), URL = ¡https://plato.stanford.edu/archives/fall2024/entries/reasons-just-vs-expl/¿.

Alvarez, Maria (2010), *Kinds of Reasons*, Oxford: Oxford University Press.

— (2018), "Reasons for Action, Acting for Reasons, and Rationality", *Synthese*, 195 (8): 3293–3310.

Archer, Sophie. (2017), 'Defending Exclusivity,' *Philosophy and Phenomenological Research*, 94: 326-341.

Audi, Robert. (2019), 'Understanding, Self-Evidence, and Justification,' *Philosophy and Phenomenological Research*, 99: 358-381.

Basu, Rima. and Schroeder, Mark. (2019), 'Doxastic Wronging,' in Kim, B. (ed) and McGrath, M. (ed) *Pragmatic Encroachment in Epistemology*, pp. 181-205. (eds.) Routledge.

Bennett, Jonathan. (1990), 'Why is Belief Involuntary?,' *Analysis*, 100.

Bergmann, Michael. (2000), 'Deontology and Defeat,' *Philosophy and Phenomenological Research*, 60: 87-102.

Bird, Allexander. (2007), 'Justified Judging', *Philosophy and Phenomenological Research*, 74: 81-110.

Bittner, Rüdiger (2001), *Doing Things for Reasons*, Oxford: Oxford University Press.

Cauterucci, Christina. (2018), 'Christine Blasey Ford changed everything,' *Slate Magazine*. Retrieved April 28, 2022, from https://slate.com/human-interest/2018/10/christine-blasey-ford-brett-kavanaugh-metoo-toxic-masculinity.html

Chrisman, Matthew. (2018), 'Epistemic Normativity and Cognitive Agency,' *Nous*, 52: 508-29.

Comesaña, Juan & McGrath, Matthew (2014), "Having False Reasons", in: *Epistemic Norms*, ed. by John Turri & Clayton Littlejohn, Oxford: Oxford University Press: 59–80.

Cunningham, J.J. (2021), 'The Matter of Motivating Reasons,' *Philosophical Studies*, 179: 1563-1589.

Dancy, Jonathan (1977), "The Logical Conscience", *Analysis*, 37 (2): 81–84.

— (2000), *Practical Reality*, Oxford: Oxford University Press.

— (2003), "Précis of Practical Reality", *Philosophy and Phenomenological Research*, 67 (2): 423–428.

— (2009), "Reasons and Rationality", in: *Spheres of Reason: New Essays in the Philosophy of Normativity*, ed. by Simon Robertson, Oxford: Oxford University Press.

— (2011), "Acting in Ignorance", *Frontiers of Philosophy in China*, 6 (3): 345–357.

– (2018). *Practical Shape: A Theory of Practical Reasoning*. Oxford: Oxford University Press.
Davidson, D. (1963), "Actions, Reasons, and Causes", *Journal of Philosophy*, 60 (23): 3–19, Reprinted in Davidson (2001).
— (2001), *Essays on Actions and Events*, Oxford: Oxford University Press.
Dular, Nicole. & Fortier, Nikki. 'Epistemic Reasons, Transparency, and Evolutionary Debunking,' *Philosophia* 49: 1455-73.
Engel-Hawbecker, N.L. (forthcoming), "Must Your Reasons Move You?," *Philosophical Studies*.
Foley, Richard. (1993), *Working Without a Net*. Oxford University Press, p. 16.
Gao, Jie. (2019), 'Credal Pragmatism,' *Philosophical Studies*, 176: 1595-1617.
Gardiner, Georgi. (2018), 'Evidentialism and Moral Encroachment,' in McCain, K. (ed), *Believing in Accordance with the Evidence: New Essays on Evidentialism*. Dordrecht: Springer.
Gibbons, John (2010). *Things That Make Things Reasonable*, Philosophy and Phenomenological Research 81 (2):335-361.
Hieronymi, Pamela. (2006), 'Controlling Attitudes,' *Pacific Philosophical Quarterly*, 87: 45-74.
Hyman, John (1999), "How Knowledge Works", *Philosophical Quarterly*, 50 (197): 433–451.
— (2010), "The Road to Larissa", *Ratio*, 23 (4): 393–414.
— (2015), *Action, Knowledge, and Will*, Oxford: Oxford University Press.
Kazemi, Alireza. (2022), 'Transparency and the Truth Norm of Belief,' *Synthese* 200: 1-18.
Kiesewetter, Benjamin. (2022), 'Are Epistemic Reasons Normative?,' *Nous* 56 (3): 670-695.
Littlejohn, Clayton (2012), *Justification and the Truth-Connection*, Cambridge: Cambridge University Press.
Lord, Errol and Sylvan, Kurt. (2019), 'Prime Time (for the Basing Relation).' In J. Adam Carter (Ed.) and Patrick Bondy (Ed.) *Well-founded Belief* (pp.141-173). London and New York: Routledge.
Marcus, Eric (2012), *Rational Causation*, Camb. Mass.: Harvard Universtiy Press.
Mason, Elinor. (2003), 'Consequentialism and the 'Ought Implies Can' Principle,' *American Philosophical Quarterly*, 40: 319-31.
McHugh, Conor. (2015), 'The Illusion of Exclusivity,' *European Journal of Philosophy*, 23: 1117-1136.
Miller, Christian (2008), "Motivation in Agents", *Noûs*, 42 (2): 222–266.
Moran, Richard. (1988), 'Making Up Your Mind: Self-Interpretation and Self-Constitution,' *Ratio*, 1: 135-51.
Munoz, Daniel. (2019), 'Defeaters and Disqualifiers,' *Mind*, 128: 887-906.
Owens, David. (2003), 'Does Belief Have an Aim?,' *Philosophical Studies*, 115: 283-305.
Parfit, Derek (1997), "Reasons and Motivation", *Proceedings of the Aristotelian Society*, suppl. vol. 71 (1): 99–130.
Pollock, John. (1986), *Contemporary Theories of Knowledge*. Savage, MD: Rowman and Littlefield.

Pollock, John. and Cruz, Joseph. (1999), *Contemporary Theories of Knowledge*, 2nd Edition. Lanham, MD: Rowman and Littlefield.
Pryor, James (2007), "Reasons and That-Clauses", *Philosophical Issues*, 17 (1): 217–244.
Rinard, Susanna. (2019), 'Equal Treatment for Belief,' *Philosophical Studies*, 176: 1923-1950.
Roeber, Blake. (2020), 'Permissive Situations and Direct Doxastic Control,' *Philosophy and Phenomenological Research*, 101: 415-431.
Raz, Joseph (2005), "Instrumental Rationality: A Reprise", *Journal of Ethics and Social Philosophy*, 1 (1): 1–19.
— (2011a), From *Normativity to Responsibility*, Oxford: Oxford University Press. 29.
Raz, Joseph (2011b), "The Myth of Instrumental Rationality", in: *From Normativity to Responsibility*, ed. by Joseph Raz, Oxford: Oxford University Press.
Roessler, Johannes (2014), "Reason Explanation and the Second-Person Perspective", *Philosophical Explorations*, 3: 1–12.
Schroeder, Mark (2008), "Having Reasons", *Philosophical Studies*, 139 (1): 57–71.
Setiya, Keiran (2011), "Reasons and Causes", *European Journal of Philosophy*, 19 (1): 129–157. – (2013), "Epistemic Agency: Some Doubts", Philosophical Issues, 23 (1): 179–198.
Shah, Nishi. (2003), 'How Truth Governs Belief', *Philosophical Review*, 112: 447-82.
—(2006), 'A New Argument for Evidentialism', *The Philosophical Quarterly*, 56: 481-98.
Shah, Nishi. and Velleman, David. (2005), 'Doxastic Deliberation', *Philosophical Review*, 114: 497-534.
Sharadin, Nathaniel. (2016), 'Nothing But the Evidential Considerations?,' *Australasian Journal of Philosophy*, 94: 1-19.
Singh, Keshav (2019), "Acting and Believing Under the Guise of Normative Reasons", *Philosophy and Phenomenological Research*, 99 (2): 409–430.
Smith, Michael (1987), "The Humean Theory of Motivation", *Mind*, 96 (381): 36–61.
Steglich-Petersen, Asbjørn. (2006), 'No Norm Needed: On the Aim of Belief', *The Philosophical Quarterly*, 56: 499-516.
—(2008), 'Does Doxastic Transparency Support Evidentialism?', *dialectica*, 62: 541-7.
—(2009), 'Weighing the Aim of Belief', *Philosophical Studies*, 145: 395-405.
Stout, Rowland (2009), "Was Sally's Reason for Running from the Bear that She Thought it was Chasing Her?", in: *New Essays in the Explanation of Action*, ed. by Constantine Sandis, Palgrave: Macmillan: 48–61.
Sturgeon, Scott. (2014), 'Pollock on Defeasible Reasons,' *Philosophical Studies*, 169: 105-118.
Sullivan-Bisset, Emma. (2018), 'Explaining Doxastic Transparency: Aim, Norm, or Function?,' *Synthese*, 195: 3453-3476.
—-(2017), 'Aims and Exclusivity,' *European Journal of Philosophy*, 25: 721-731.
Sullivan-Bisset, Emma. & Noordhof, Paul. (2013), 'A Defense of Owens' Exclusivity Objection to Beliefs Having Aims,' *Philosophical Studies*, 163: 453-457.
—-(2020),'The Transparent Failure of Norms to Keep Up Standards of Belief,' *Philosophical Studies*, 177: 1213-27.

Sylvan, Kurt. (2016), 'The Illusion of Discretion,' *Synthese*, 193: 1635-1665.
Turri, John (2009), "The Ontology of Epistemic Reasons", *Noûs*, 43 (3): 490–512.
Vahid, Hamid. (2022), "No Practical Reasons for Belief: the Epistemic Significance of Practical Considerations," *Synthese* 200 (2): 1-18.
Velleman, David. (2000), 'On the Aim of belief.' In D. Velleman (Ed.), *The Possibility of Practical Reason*, (pp. 244–281). Oxford: Oxford University Press.
Wallace, R. Jay (2003), "Explanation, Deliberation, and Reasons", *Philosophy and Phenomenaological Research*, 67 (2): 429–435.
Way, Jonathan. (2016), 'Two Arguments for Evidentialism,' *Philosophical Quarterly*, 66: 805-818.
—(2016), 'If You Justifiably Believe that You Ought to Φ, You Ought to Φ,' *Philosophical Studies*, 173: 1873-1895.
Winters, Barbara. (1979), 'Believing at Will,' *Journal of Philosophy*, 76.
Williams, Bernard. (1973), 'Deciding to Believe', in: *Problems of the Self*. Cambridge, UK: Cambridge University Press.